

Arjun Mehta

AI/ML Engineer | Generative AI | LLMOps | Distributed Systems
Email: arjun.mehta.ai@gmail.com | Phone: +91-9876543210
LinkedIn: linkedin.com/in/arjunmehta-ai | GitHub: github.com/arjunmehta-ai | Location: Bangalore, India

Professional Summary

AI/ML Engineer with 4+ years of experience designing, developing, and deploying machine learning systems and Generative AI solutions at scale. Experienced in building LLM-powered applications, distributed data pipelines, and production-grade ML systems. Strong expertise in Python, PyTorch, TensorFlow, Spark, Kafka, and cloud platforms including AWS and Azure. Proven ability to optimize model performance, implement LLMOps pipelines, and deploy scalable AI solutions supporting thousands of users.

Technical Skills

Programming	Python, Java, Scala, SQL
Machine Learning	PyTorch, TensorFlow, scikit-learn, NLP, Deep Learning, Model Evaluation
Generative AI	LLMs, Prompt Engineering, RAG, Fine-tuning, LLMOps
Big Data	Apache Spark, Hadoop, Kafka, HBase
Cloud & DevOps	AWS, Azure, Docker, Kubernetes, CI/CD
Databases	PostgreSQL, Redis, Aerospike, NoSQL

Professional Experience

Senior Machine Learning Engineer – FinTech AI Labs (2022 – Present)

- Developed and fine-tuned Generative AI models and LLM-powered personalization systems serving 500k+ users.
- Designed scalable distributed ML pipelines using Apache Spark and Kafka for real-time event processing.
- Implemented LLMOps workflows for training, evaluation, monitoring, and deployment of large language models.
- Deployed containerized ML services using Docker and Kubernetes on AWS and Azure cloud platforms.
- Improved model inference latency by 35% through optimization and model quantization.

Machine Learning Engineer – DataScale Technologies (2020 – 2022)

- Built and deployed machine learning models for fraud detection and personalization systems.
- Developed batch and real-time data pipelines using Hadoop, Spark, and Kafka.
- Implemented model monitoring, performance evaluation, and automated retraining pipelines.
- Collaborated with cross-functional product and engineering teams to deploy AI products in production.

Key AI/ML Projects

LLM Customer Support Assistant

- Built a Retrieval-Augmented Generation (RAG) chatbot using PyTorch, LangChain, and vector databases.
- Reduced customer support response time by 40% and improved response accuracy to 92%.

Real-Time Fraud Detection System

- Designed a streaming ML system using Kafka and Spark Streaming to detect fraud in real-time transactions.
- Achieved 96% detection accuracy with sub-second prediction latency.

Education

M.Tech in Artificial Intelligence – Indian Institute of Technology (IIT) Delhi (2020)

B.Tech in Computer Science – National Institute of Technology (NIT) Trichy (2018)

Certifications

- AWS Certified Machine Learning – Specialty
- Deep Learning Specialization – DeepLearning.AI
- Machine Learning Engineering for Production (MLOps) – Coursera